

Comparative Analysis of Trie, AVL, Segment, and Suffix Trees for NLP Based Searching and Text Analysis

Team Members Name:

K. Gnandee (106124057)

T. Aditya (106124011)

Devarakonda Harshavardhanachari (106124031)

Submitted to:

Dr. B. Nithya Mam

Course:

Data Structures[CSPC32]

Institution:

National institute of technology, Trichy

Abstraction

More usage of textual data in every app, webpages has made the efficient search of text data as an important research domain in Natural Language Processing (NLP). So as data is stored in the form of text everywhere we are trying many efficient methods such that it is helpful to retrieve and store data to user. Linguistic analysis of this data is also more important here. Traditional data structures used for searching and indexing face significant challenges when adapted to semantic, contextual, and large-scale textual data.

Our project explores a comparative implementation and analysis of four advanced tree-based data structures—**Trie**, **AVL Tree**, **Segment Tree**, and **Suffix Tree**—in the context of NLP-driven text retrieval and analysis tasks.

Each structure is analysed based on its ability to perform operations like prefix searching, ordered retrieval, range-based sentiment aggregation, and substring detection. The outcomes are checked for all types of trees. There is a comparison in all off the trees with respect to the trade-offs between time complexity, space complexity, and linguistic applicability. The findings help us in selecting the most suitable data structure for specific NLP tasks, bridging traditional data structures with modern text analytics.

Introduction

Natural Language Processing (NLP) is a computational discipline that enables machines to interpret and process human language. As NLP systems scale to handle huge amount of textual data, selecting efficient data structures becomes crucial for maintaining good results according to the user and contextual understanding. While in earlier days text-processing systems primarily used simple search or hashing mechanisms, modern NLP applications—such as autocompletion, sentiment tracking, and plagiarism detection—demand

specialized data structures that can efficiently store and query linguistically meaningful data.

In this project, four significant tree structures—Trie, AVL, Segment, and Suffix Trees—are examined for their performance and applicability to NLP tasks. Each structure has its unique strengths: Tries can do the prefix-based searches faster, AVL Trees are efficient in ordered lookups, Segment Trees are efficient in range-based analysis, and Suffix Trees are efficient in substring matching and text similarity detection.

The initial objective of our project is to measure **time efficiency**, **memory utilization**, and **scalability** of these structures within the NLP context, thereby identifying which structure best suits a given language-processing scenario.

Literature Background

Tree-based data structure are used in algorithmic text processing.

- **Trie Trees** were first proposed as an efficient way to store dynamic sets of strings with shared prefixes.
- **AVL Trees**, named after Adelson-Velsky and Landis, introduced self-balancing properties that ensure consistent logarithmic access time for ordered data.
- **Segment Trees** emerged in range-query optimization, allowing quick computation of aggregated metrics over sub-ranges—an operation that maps elegantly to sentiment and topic scoring in NLP.
- **Suffix Trees**, originating from string pattern matching algorithms, provide linear-time substring detection and are foundational in bioinformatics and text similarity analysis.

Our project unifies them under a single NLP pipeline and evaluates their relative strengths, limitations, and computational efficiency.

Methodology

In our project we are implementing all of four data structures and integrates them into a text-processing pipeline designed for linguistic operations. Each implementation are efficient in specific NLP functions, evaluated for speed, scalability, and memory efficiency.

Implementation Flow

The NLP system operates in parallel data-structure pipelines:

1. Input text is tokenized and passed through each data structure.
2. Trie handles prefix lookups and suggestion generation.
3. AVL Tree organizes and ranks tokens based on semantics.
4. Segment Tree aggregates sentiment or keyword frequencies across text segments.
5. Suffix Tree processes substring matching and similarity detection.
6. Outputs from all modules are integrated into an analysis dashboard for visualization and comparison.

Trie Tree Implementation

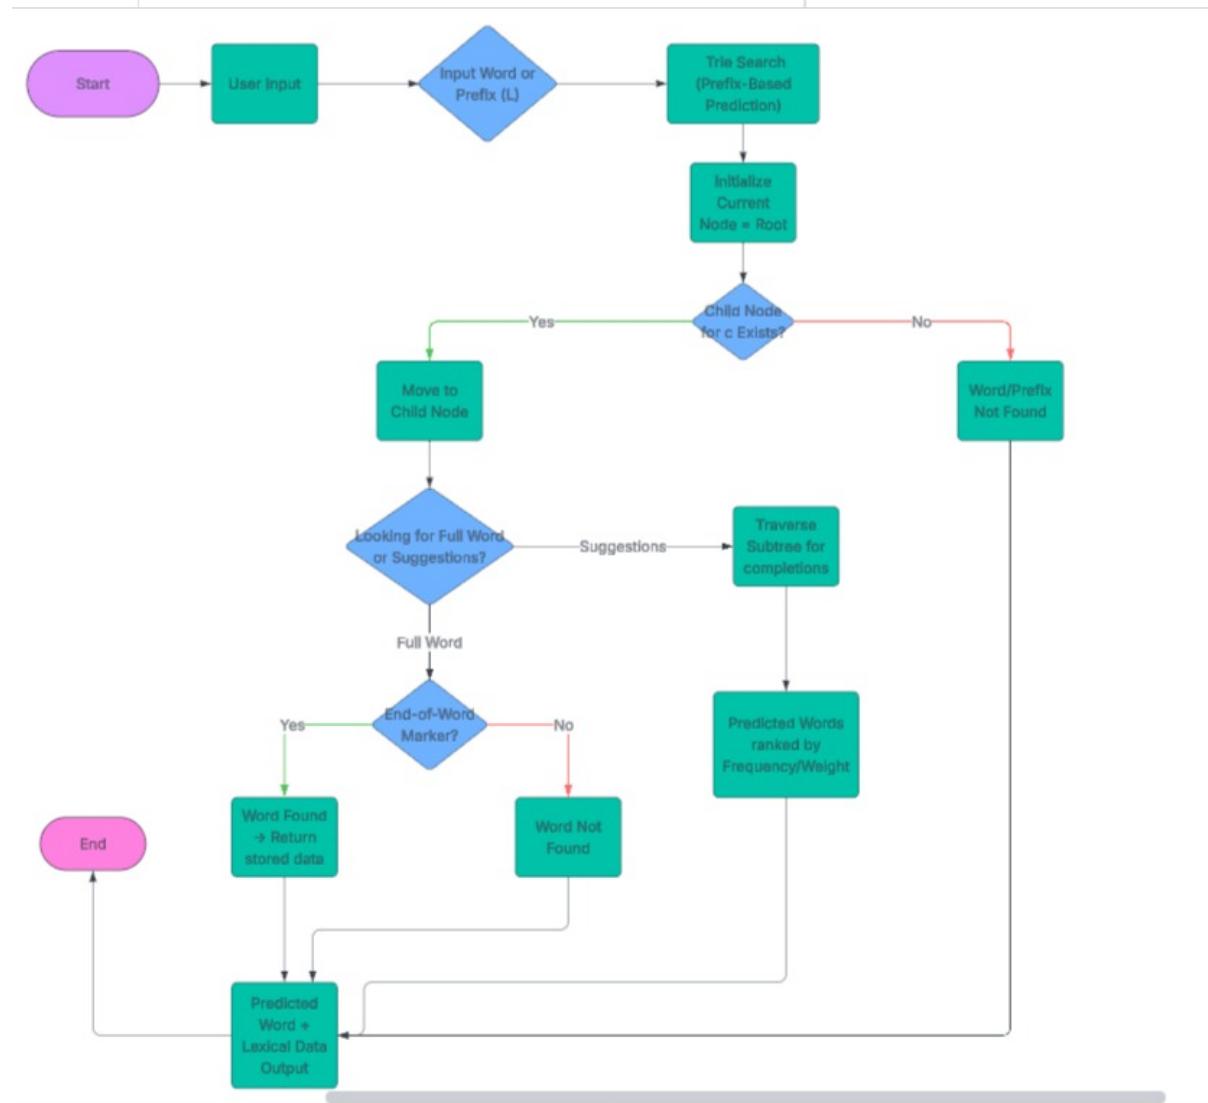
The **Trie**, or prefix tree, was implemented as the backbone for the **autocomplete** and **lexical prediction** module. Each node represents a single character, forming paths that represent complete words.

Enhancements included:

- **Frequency weighting**, allowing the system to prioritize more common words.
- **Synonym linking**, enabling semantic completion rather than strict prefix matching.
- **Context-awareness**, where suggestions adapt based on previously used tokens.

This approach ensures $O(L)$ lookup time (where L is word length) and enables rapid suggestion generation, making the Trie ideal for real-time NLP applications such as search bars and dialogue systems.

Trie Implementation Flowchart



This flowchart shows how text (data) are added and searched in the Trie data structure. Each letter is stored as a node, and when a word ends, it's marked (like \0 in string). It helps quickly find all words starting with a given prefix.

AVL Tree Implementation

The **AVL Tree**, a self-balancing binary search tree, ensures that height differences between subtrees are minimal, maintaining $O(\log N)$ complexity for insertions and lookups. The balance factor(bf) for each lies in the range(-2<bf<2) where bf is an integer .

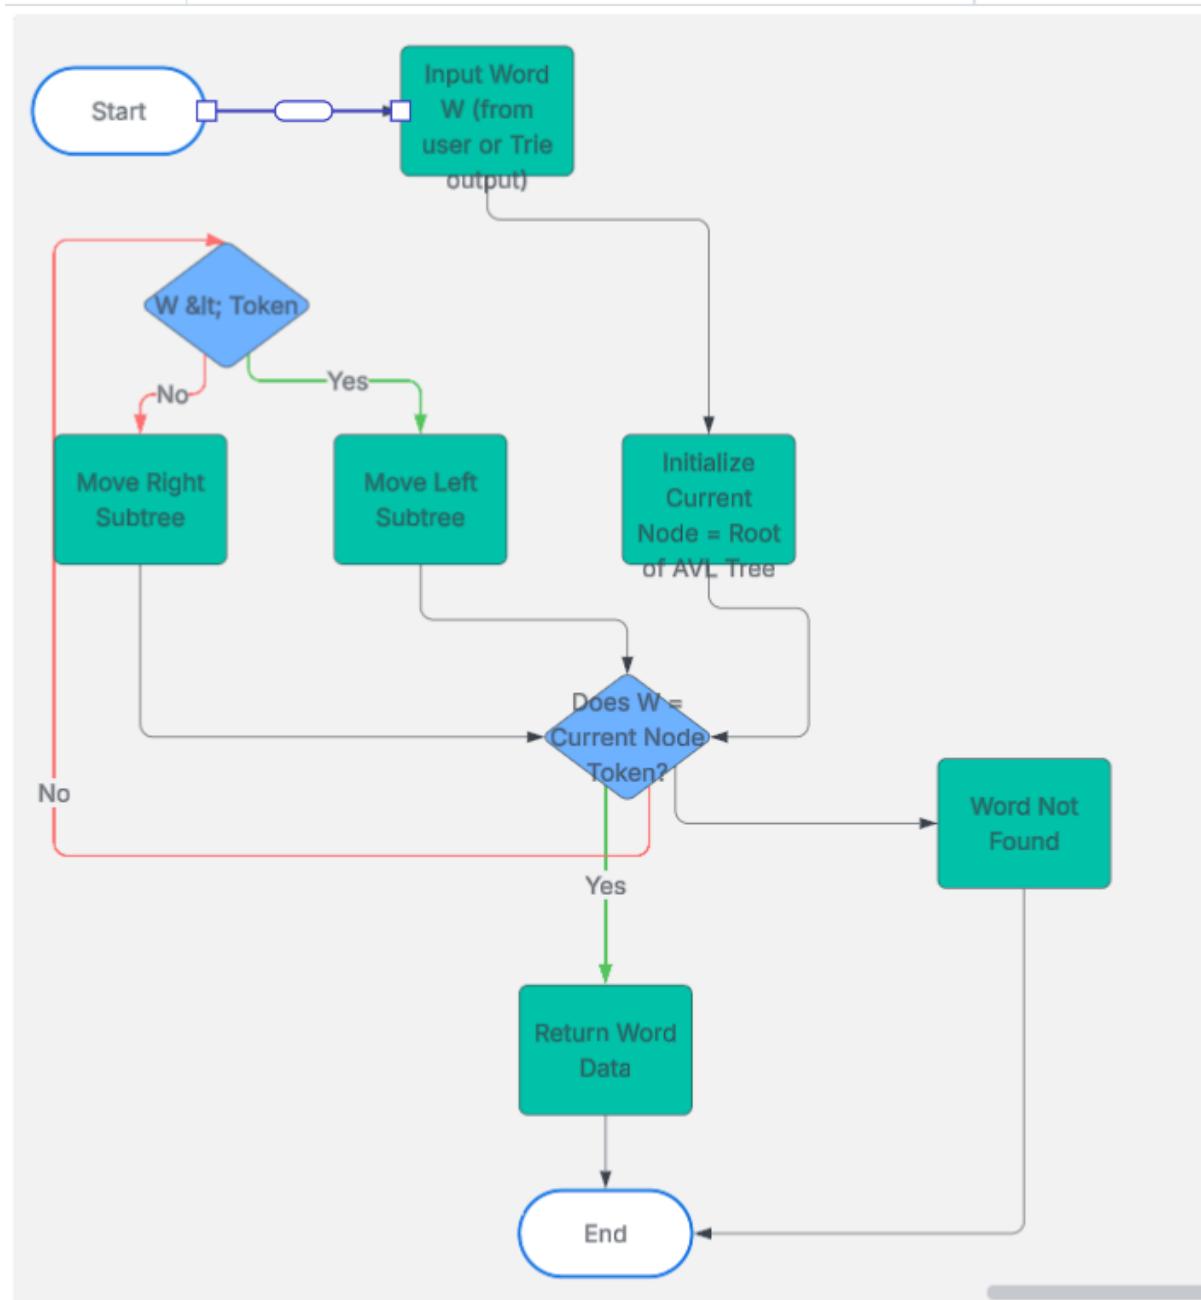
In our project, the AVL Tree was used to build an **ordered lexical database** where each node stores:

- Word token
- Part-of-Speech (POS) tag
- Semantic category or identifier

Such organization supports queries like:

- Retrieve all verbs ranked by frequency
- Fetch the top-N nouns contributing to positive sentiment

Its ordered nature makes AVL Trees especially effective for dictionary-style queries, semantic ranking, and frequency-based retrievals, though the balancing overhead increases memory usage slightly.



AVL Tree Implementation Flowchart

This flowchart explains how the AVL Tree keeps data balanced while inserting or searching. It adjusts itself using rotations so that searching and sorting of words stay fast and efficient.

Segment Tree Implementation

The **Segment Tree** excels in range-based queries. For NLP, the text corpus was divided into logical segments such as sentences or paragraphs. Each node maintains aggregate statistics like:

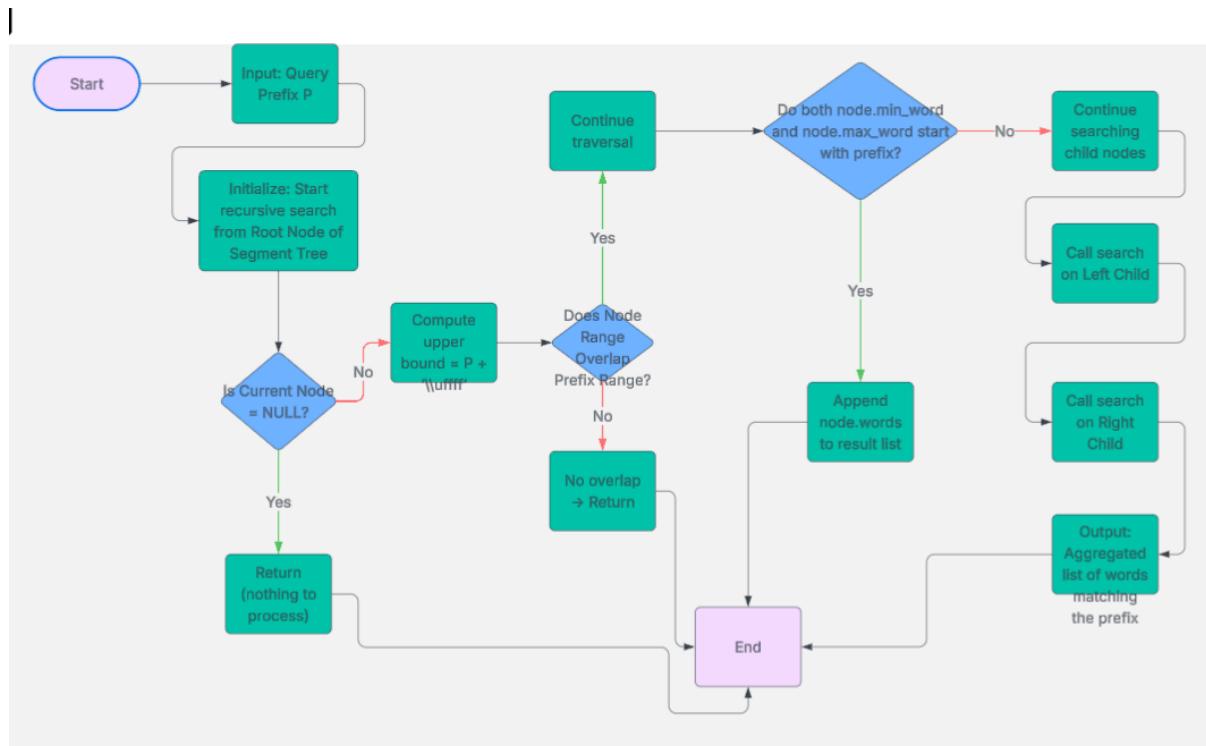
- Sentiment polarity (positive/negative/neutral)
- Keyword frequency or topic density

The Segment Tree allows efficient range queries such as:

- “What is the average sentiment between chapters 3 and 5?”
- “Which section has the highest topic concentration?”

With $O(\log N)$ query and update complexity, it is ideal for **trend tracking**, **sentiment progression analysis**, and **topic segmentation** across lengthy documents.

Segment Tree Implementation Flowchart



This flowchart shows how the text(data) is divided into parts and how the segment tree stores the values like keyword counts. It helps to find or update data in any range quickly.

Suffix Tree Implementation

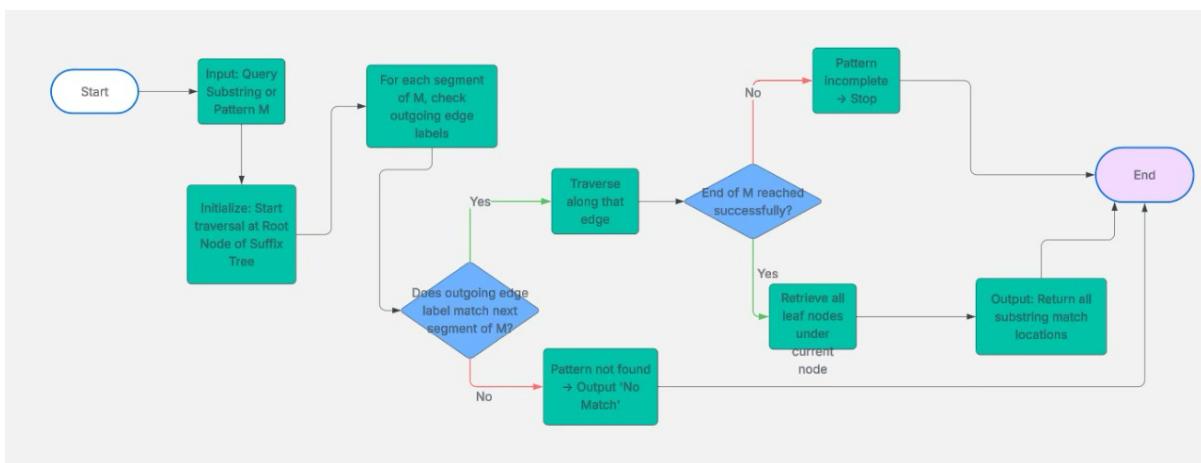
The **Suffix Tree** provides a powerful structure for substring and pattern-matching tasks. Constructed in $O(N)$ time, it enables searches for substrings or repeated patterns in $O(M)$ time (M = length of query substring).

Applications implemented:

- **Plagiarism and duplication detection** across multiple documents.
- **Phrase repetition and pattern identification.**
- **Contextual overlap** discovery for semantic similarity detection.

The Suffix Tree's unmatched substring search performance makes it invaluable in NLP applications requiring fine-grained text analysis. Suffix Tree is memory intensive used tree data type.

Suffix Tree Implementation Flowchart

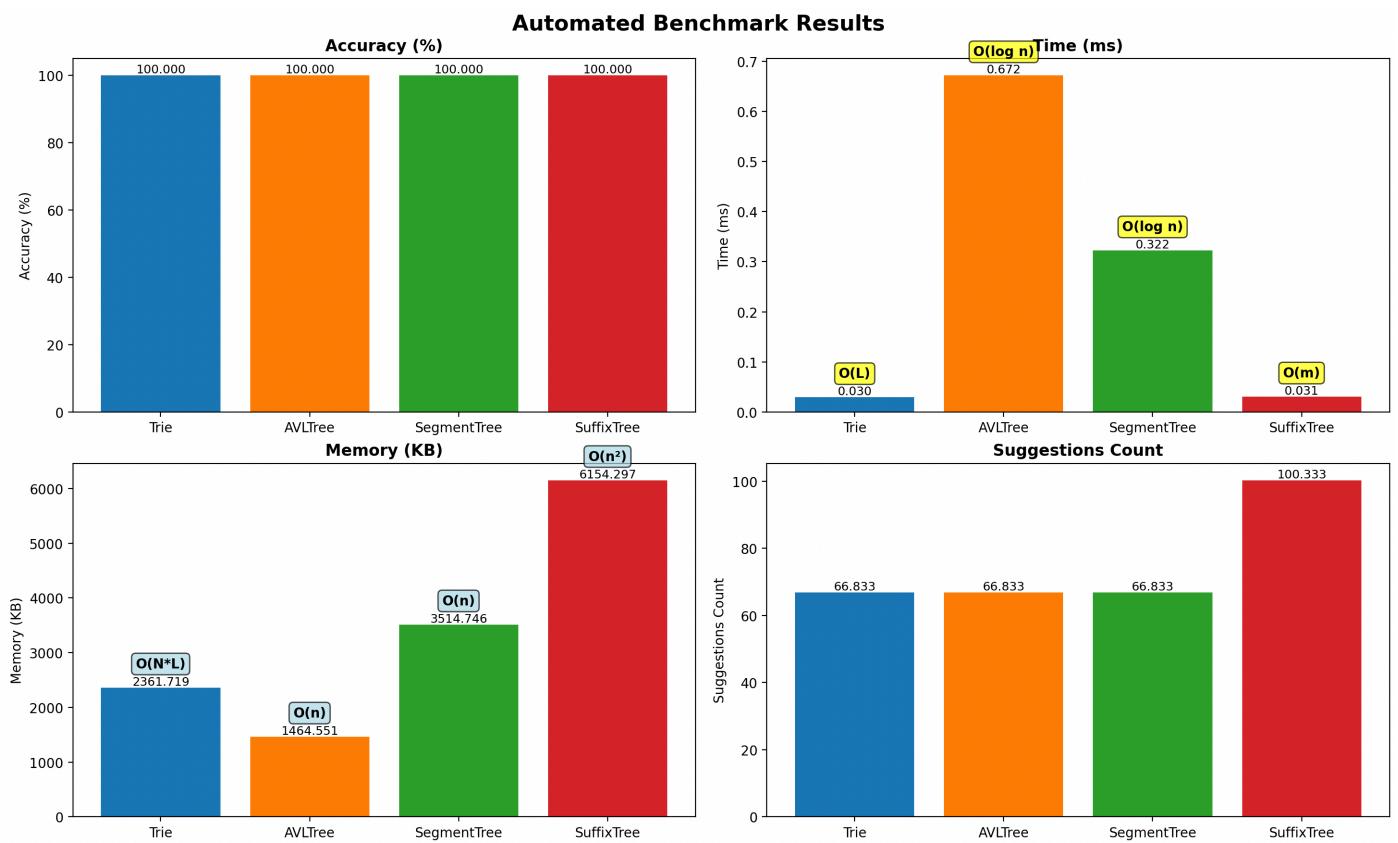


This flowchart explains how all possible endings (suffixes) of a text(data) are stored in a suffix tree. It is used to find repeating patterns or substrings in large text easily and fast.

Comparative Analysis

The four structures were analysed based on their performance in four primary NLP tasks: prefix searching, ordered retrieval, range sentiment evaluation, and substring matching.

Feature / Use Case	Trie	AVL Tree	Segment Tree	Suffix Tree
Prefix Autocomplete	Excellent	Good	Moderate	Weak
Substring / Pattern Matching	Excellent	Weak	Weak	Excellent
Memory Usage	Moderate	High	Excellent	Moderate



Analysis Discussion:

- **Trie Trees** dominate prefix-based operations due to direct traversal without comparisons.
- **AVL Trees** are superior when ordered search or semantic sorting is required.
- **Segment Trees** handle range aggregation efficiently, supporting analytical NLP tasks.
- **Suffix Trees** outperform others in substring and pattern detection, crucial for plagiarism and overlap detection.

The trade-off lies in memory consumption: Trie and Suffix Trees require significant space, while AVL and Segment Trees are more compact.

Results

Performance Metrics

- **Time Complexity:**
Trie and Suffix Trees perform string-based operations in $O(L)$ and $O(M)$, while AVL and Segment Trees maintain $O(\log N)$ efficiency for numeric and ordered operations.
- **Space Efficiency:**
Trie and Suffix Trees consume more memory due to node-rich storage; AVL and Segment Trees are more space-efficient.
- **Scalability:**
Segment Trees scale gracefully for large documents; Trie and Suffix Trees need optimization for sparse datasets.

Functional Outputs

- **Autocomplete Suggestions** from the Trie.

- **Sentiment Trend Visualization** from Segment Tree queries.
- **Plagiarism Detection and Overlap Analysis** using Suffix Trees.

Analytical Observations

Smaller database tests show that Trie achieves the lowest lookup time for prefix tasks, while Segment Tree provides the most stable performance over large text ranges. Suffix Tree, despite its space cost, delivers unmatched substring detection capabilities.

Conclusion

The comparative analysis confirms that no single data structure outperforms others(cannot overcome other trees efficiencies) .

- The **Trie** remains the most efficient for prefix-based prediction and intelligent autocomplete.
- The **AVL Tree** offers superior performance in structured, ordered retrieval operations.
- The **Segment Tree** dominates analytical and range-based computations.
- The **Suffix Tree** proves essential for substring and similarity detection tasks.

While Trie and Suffix Trees are more memory-intensive, their precision and contextual intelligence justify the cost for NLP applications. Future work will explore integrating neural embeddings (word vectors) into these trees, forming hybrid tree-graph architectures capable of semantic-level querying beyond lexical boundaries.

References

1. **Cormen, Leiserson, Rivest, Stein — *Introduction to Algorithms (MIT Press)*:** Foundational text for all discussed data structures.
2. **Jurafsky & Martin — *Speech and Language Processing*:** Core reference for NLP algorithmic approaches.
3. **Dan Gusfield — *Algorithms on Strings, Trees, and Sequences*:** Key work on suffix and string-matching structures.
4. **Hachiya, K. et al. (IEEE):** “Efficient Trie-Based Query Systems for Natural Language Retrieval.”
5. **GeeksforGeeks — “Comparison of Trie, AVL, Segment, and Suffix Trees”:** Practical examples and code explanations.
6. **Stanford NLP Group Publications:** Studies integrating linguistic data with algorithmic structures.
7. **Karp, R.M. (ACM Journal) — “Pattern Matching Algorithms and Applications.”**
8. **Coursera Course — “Advanced Data Structures in C++” by UC San Diego:** Includes Trie, Segment, and Suffix Tree implementations.
9. **IEEE Xplore Paper — “Range-Query Optimization for Text Mining using Segment Trees.”**
10. **GitHub Repository — “Trie-AVL-SuffixTree-NLP” (open-source comparative analysis):** Implements multi-tree text search pipeline for NLP experiments.